

Project 3 : Supervised Machine Learning - Classification

Objective

The main objective of my analysis is to focus on prediction that the person would be dead or alive based on the covid and health conditions given in the dataset.

Dataset

The dataset was provided by the Mexican government ([link](#)). This dataset contains an enormous number of anonymized patient-related information including pre-conditions. The raw dataset consists of 21 unique features and 1,048,576 unique patients. In the Boolean features, 1 means "yes" and 2 means "no". values as 97 and 99 are missing data.

sex: 1 for female and 2 for male. age: of the patient. classification: covid test findings. Values 1-3 mean that the patient was diagnosed with covid in different degrees. 4 or higher means that the patient is not a carrier of covid or that the test is inconclusive. patient type: type of care the patient received in the unit. 1 for returned home and 2 for hospitalization. pneumonia: whether the patient already have air sacs inflammation or not. pregnancy: whether the patient is pregnant or not. diabetes: whether the patient has diabetes or not. copd: Indicates whether the patient has Chronic obstructive pulmonary disease or not. asthma: whether the patient has asthma or not. inmsupr: whether the patient is immunosuppressed or not. hypertension: whether the patient has hypertension or not. cardiovascular: whether the patient has heart or blood vessels related disease. renal chronic: whether the patient has chronic renal disease or not. other disease: whether the patient has other disease or not. obesity: whether the patient is obese or not. tobacco: whether the patient is a tobacco user. usmer: Indicates whether the patient treated medical units of the first, second or third level. medical unit: type of institution of the National Health System that provided the care. intubed: whether the patient was connected to the ventilator. icu: Indicates whether the patient had been admitted to an Intensive Care Unit. date died: If the patient died indicate the date of death, and 9999-99-99 otherwise.

Dependant Variable :

alive : it will be created during EDA through Date_died, Whether Patient died or not.

Import Packages

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter

from sklearn.preprocessing import OrdinalEncoder
```

```

from sklearn.feature_selection import chi2, SelectFromModel,
SelectKBest
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import classification_report, roc_auc_score,
ConfusionMatrixDisplay, confusion_matrix,
precision_recall_fscore_support
from sklearn.model_selection import GridSearchCV
from imblearn.over_sampling import SMOTE, ADASYN

```

#Load the dataset

```

data= pd.read_csv("Covid Data.csv")
data

```

	USMER	MEDICAL_UNIT	SEX	PATIENT_TYPE	DATE_DIED	
INTUBED \						
0	2	1	1	1	03/05/2020	97
1	2	1	2	1	03/06/2020	97
2	2	1	2	2	09/06/2020	1
3	2	1	1	1	12/06/2020	97
4	2	1	2	1	21/06/2020	97
...
1048570	2	13	2	1	9999-99-99	97
1048571	1	13	2	2	9999-99-99	2
1048572	2	13	2	1	9999-99-99	97
1048573	2	13	2	1	9999-99-99	97
1048574	2	13	2	1	9999-99-99	97

	PNEUMONIA	AGE	PREGNANT	DIABETES	...	ASTHMA	INMSUPR	\
0	1	65	2	2	...	2	2	
1	1	72	97	2	...	2	2	
2	2	55	97	1	...	2	2	
3	2	53	2	2	...	2	2	
4	2	68	97	1	...	2	2	
...	
1048570	2	40	97	2	...	2	2	
1048571	2	51	97	2	...	2	2	

1048572	2	55	97	2	...	2	2
1048573	2	28	97	2	...	2	2
1048574	2	52	97	2	...	2	2

	HIPERTENSION	OTHER_DISEASE	CARDIOVASCULAR	OBESITY
RENAL_CHRONIC \				
0	1	2	2	2
2				
1	1	2	2	1
1				
2	2	2	2	2
2				
3	2	2	2	2
2				
4	1	2	2	2
2				
...
...				
1048570	2	2	2	2
2				
1048571	1	2	2	2
2				
1048572	2	2	2	2
2				
1048573	2	2	2	2
2				
1048574	2	2	2	2
2				

	TOBACCO	CLASIFFICATION_FINAL	ICU
0	2	3	97
1	2	5	97
2	2	3	2
3	2	7	97
4	2	3	97
...
1048570	2	7	97
1048571	2	7	2
1048572	2	7	97
1048573	2	7	97
1048574	2	7	97

[1048575 rows x 21 columns]

EDA

EDA consists of:

- Look for missing values, we already know 97-99 are the missing value

- Fill or remove missing values, here we have 1st replaced the missing values of pregnant for man as '2', ICU and INTUBED as '2' where PATIENT_TYPE was 2 as patient were sent home, rest were dropped.
- See the value count of each column to get details of distribution.
- Add a new dependable classification feature 'Alive', such that where DEATH_DATE is '9999-99-99' the person is alive i.e '0' and others mark as not alive i.e '1'
- Make a heatmap to get a better understanding of relationship between dependant and independent features. As it shows the correlation between each feature.
- Create ordinal encoding for age as it would create better relation with dependant variable.
- Using Count plot of different feature with hue as 'Alive' to get a clear understanding of distribution throug visualization.
- Drop the unnecessary column such as 'AGE', 'DEATH_DATE' as it will create unnecessary noise.
- Get feature importance through chi as data is unbalanced.

Observation

- There are no null values in any column
- The dataset is very unbalanced as almost every column has a majority class with larger imbalance ratio.

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1048575 entries, 0 to 1048574
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   USMER                                1048575 non-null  int64
1   MEDICAL_UNIT                        1048575 non-null  int64
2   SEX                                  1048575 non-null  int64
3   PATIENT_TYPE                        1048575 non-null  int64
4   DATE_DIED                           1048575 non-null  object
5   INTUBED                             1048575 non-null  int64
6   PNEUMONIA                          1048575 non-null  int64
7   AGE                                  1048575 non-null  int64
8   PREGNANT                            1048575 non-null  int64
9   DIABETES                            1048575 non-null  int64
10  COPD                                1048575 non-null  int64
11  ASTHMA                              1048575 non-null  int64
12  INMSUPR                             1048575 non-null  int64
13  HIPERTENSION                        1048575 non-null  int64
14  OTHER_DISEASE                       1048575 non-null  int64
15  CARDIOVASCULAR                     1048575 non-null  int64
16  OBESITY                             1048575 non-null  int64
17  RENAL_CHRONIC                      1048575 non-null  int64
18  TOBACCO                             1048575 non-null  int64
19  CLASIFFICATION_FINAL               1048575 non-null  int64
```

```
20 ICU 1048575 non-null int64
dtypes: int64(20), object(1)
memory usage: 168.0+ MB
```

```
for column in data.columns:
    print(data[column].value_counts(), end="\n\n")
```

```
USMER
2 662903
1 385672
Name: count, dtype: int64
```

```
MEDICAL_UNIT
12 602995
4 314405
6 40584
9 38116
3 19175
8 10399
10 7873
5 7244
11 5577
13 996
7 891
2 169
1 151
Name: count, dtype: int64
```

```
SEX
1 525064
2 523511
Name: count, dtype: int64
```

```
PATIENT_TYPE
1 848544
2 200031
Name: count, dtype: int64
```

```
DATE_DIED
9999-99-99 971633
06/07/2020 1000
07/07/2020 996
13/07/2020 990
16/06/2020 979
...
24/11/2020 1
17/12/2020 1
08/12/2020 1
16/03/2021 1
22/04/2021 1
```

Name: count, Length: 401, dtype: int64

INTUBED

97 848544

2 159050

1 33656

99 7325

Name: count, dtype: int64

PNEUMONIA

2 892534

1 140038

99 16003

Name: count, dtype: int64

AGE

30 27010

31 25927

28 25313

29 25134

34 24872

...

114 2

116 2

111 1

121 1

113 1

Name: count, Length: 121, dtype: int64

PREGNANT

97 523511

2 513179

1 8131

98 3754

Name: count, dtype: int64

DIABETES

2 920248

1 124989

98 3338

Name: count, dtype: int64

COPD

2 1030510

1 15062

98 3003

Name: count, dtype: int64

ASTHMA

2 1014024

```
1      31572
98      2979
Name: count, dtype: int64
```

```
INMSUPR
2      1031001
1      14170
98      3404
Name: count, dtype: int64
```

```
HIPERTENSION
2      882742
1      162729
98      3104
Name: count, dtype: int64
```

```
OTHER_DISEASE
2      1015490
1      28040
98      5045
Name: count, dtype: int64
```

```
CARDIOVASCULAR
2      1024730
1      20769
98      3076
Name: count, dtype: int64
```

```
OBESITY
2      885727
1      159816
98      3032
Name: count, dtype: int64
```

```
RENAL_CHRONIC
2      1026665
1      18904
98      3006
Name: count, dtype: int64
```

```
TOBACCO
2      960979
1      84376
98      3220
Name: count, dtype: int64
```

```
CLASIFFICATION_FINAL
7      499250
3      381527
6      128133
```

```

5      26091
1      8601
4      3122
2      1851
Name: count, dtype: int64

```

```

ICU
97      848544
2      175685
1      16858
99      7488
Name: count, dtype: int64

```

```

data['Alive'] = data['DATE_DIED'].apply(lambda x: 1 if x == '9999-99-99' else 2)

```

```

data['Alive'].value_counts()

```

```

Alive
1      971633
2      76942
Name: count, dtype: int64

```

```

mask = data['PREGNANT'].isin([97, 98, 99])
data['PREGNANT'] = data['PREGNANT'].mask(mask & (data['SEX'] == 1), None)
data['PREGNANT'] = data['PREGNANT'].mask(mask & (data['SEX'] == 2), 2)
data.loc[data['PATIENT_TYPE'] == 1, ['INTUBED', 'ICU']] = 2
refined_data = data.replace(to_replace=[97,98,99], value=None)
refined_data['AGE'] = data['AGE']

```

```

refined_data.dropna(inplace=True)
refined_data

```

	USMER	MEDICAL_UNIT	SEX	PATIENT_TYPE	DATE_DIED	INTUBED
PNEUMONIA \						
0	2	1	1	1	03/05/2020	2
1						
1	2	1	2	1	03/06/2020	2
1						
2	2	1	2	2	09/06/2020	1
2						
3	2	1	1	1	12/06/2020	2
2						
4	2	1	2	1	21/06/2020	2
2						
...
...						
1048570	2	13	2	1	9999-99-99	2
2						

1048571	1	13	2	2	9999-99-99	2
2						
1048572	2	13	2	1	9999-99-99	2
2						
1048573	2	13	2	1	9999-99-99	2
2						
1048574	2	13	2	1	9999-99-99	2
2						

	AGE	PREGNANT	DIABETES	...	INMSUPR	HIPERTENSION
OTHER_DISEASE \						
0	65	2.0	2	...	2	1
2						
1	72	2.0	2	...	2	1
2						
2	55	2.0	1	...	2	2
2						
3	53	2.0	2	...	2	2
2						
4	68	2.0	1	...	2	1
2						
...

1048570	40	2.0	2	...	2	2
2						
1048571	51	2.0	2	...	2	1
2						
1048572	55	2.0	2	...	2	2
2						
1048573	28	2.0	2	...	2	2
2						
1048574	52	2.0	2	...	2	2
2						

	CARDIOVASCULAR	OBESITY	RENAL_CHRONIC	TOBACCO
CLASIFFICATION_FINAL ICU \				
0	2	2	2	2
3 2				
1	2	1	1	2
5 2				
2	2	2	2	2
3 2				
3	2	2	2	2
7 2				
4	2	2	2	2
3 2				
...
...
1048570	2	2	2	2

7	2				
1048571		2	2	2	2
7	2				
1048572		2	2	2	2
7	2				
1048573		2	2	2	2
7	2				
1048574		2	2	2	2
7	2				

	Alive
0	2
1	2
2	2
3	2
4	2
...	...
1048570	1
1048571	1
1048572	1
1048573	1
1048574	1

[1019666 rows x 22 columns]

```
for column in refined_data.columns:
    print(refined_data[column].value_counts(), end="\n\n")
```

USMER

2	655076
---	--------

1	364590
---	--------

Name: count, dtype: int64

MEDICAL_UNIT

12	587800
----	--------

4	306502
---	--------

6	37675
---	-------

9	36983
---	-------

3	18566
---	-------

8	10070
---	-------

10	7511
----	------

5	7045
---	------

11	5541
----	------

7	856
---	-----

13	808
----	-----

2	158
---	-----

1	151
---	-----

Name: count, dtype: int64

SEX

```
2    510596
1    509070
Name: count, dtype: int64
```

```
PATIENT_TYPE
1    830385
2    189281
Name: count, dtype: int64
```

```
DATE_DIED
9999-99-99    946356
06/07/2020     966
07/07/2020     963
13/07/2020     947
16/06/2020     937
...
10/09/2020      1
19/04/2021      1
30/10/2020      1
24/03/2021      1
22/04/2021      1
Name: count, Length: 393, dtype: int64
```

```
INTUBED
2    986730
1    32936
Name: count, dtype: int64
```

```
PNEUMONIA
2    883254
1    136412
Name: count, dtype: int64
```

```
AGE
30    26435
31    25386
28    24708
29    24574
32    24375
...
115      2
119      2
111      1
121      1
113      1
Name: count, Length: 121, dtype: int64
```

```
PREGNANT
2.0    1011838
1.0      7828
```

Name: count, dtype: int64

DIABETES

2 898137

1 121529

Name: count, dtype: int64

COPD

2 1005398

1 14268

Name: count, dtype: int64

ASTHMA

2 989337

1 30329

Name: count, dtype: int64

INMSUPR

2 1006143

1 13523

Name: count, dtype: int64

HIPERTENSION

2 861207

1 158459

Name: count, dtype: int64

OTHER_DISEASE

2 992670

1 26996

Name: count, dtype: int64

CARDIOVASCULAR

2 999677

1 19989

Name: count, dtype: int64

OBESITY

2 863614

1 156052

Name: count, dtype: int64

RENAL_CHRONIC

2 1001446

1 18220

Name: count, dtype: int64

TOBACCO

2 937307

1 82359

```
Name: count, dtype: int64
```

```
CLASIFFICATION_FINAL
```

```
7      488131
```

```
3      375920
```

```
6      117178
```

```
5       25175
```

```
1       8397
```

```
4       3082
```

```
2       1783
```

```
Name: count, dtype: int64
```

```
ICU
```

```
2      1003258
```

```
1       16408
```

```
Name: count, dtype: int64
```

```
Alive
```

```
1      946356
```

```
2       73310
```

```
Name: count, dtype: int64
```

```
# Define age groups
```

```
bins = [0, 12, 18, 35, 50, 65, 80, 125]
```

```
labels = ['Child', 'Teen', 'Young Adult', 'Adult', 'Middle-aged',  
'Senior', 'Elderly']
```

```
# Apply binning
```

```
refined_data['Age Group'] = pd.cut(refined_data['AGE'], bins=bins,  
labels=labels, right=False)
```

```
# Ordinal Encoding
```

```
encoder = OrdinalEncoder(categories=[labels]) # Preserve order
```

```
refined_data['Age Group'] = encoder.fit_transform(refined_data[['Age  
Group']])
```

```
refined_data
```

	USMER	MEDICAL_UNIT	SEX	PATIENT_TYPE	DATE_DIED	INTUBED
PNEUMONIA \						
0	2	1	1	1	03/05/2020	2
1						
1	2	1	2	1	03/06/2020	2
1						
2	2	1	2	2	09/06/2020	1
2						
3	2	1	1	1	12/06/2020	2
2						
4	2	1	2	1	21/06/2020	2
2						

...
1048570	2	13	2	1	9999-99-99
2					
1048571	1	13	2	2	9999-99-99
2					
1048572	2	13	2	1	9999-99-99
2					
1048573	2	13	2	1	9999-99-99
2					
1048574	2	13	2	1	9999-99-99
2					

	AGE	PREGNANT	DIABETES	...	HIPERTENSION	OTHER_DISEASE	\
0	65	2.0	2	...	1		2
1	72	2.0	2	...	1		2
2	55	2.0	1	...	2		2
3	53	2.0	2	...	2		2
4	68	2.0	1	...	1		2
...
1048570	40	2.0	2	...	2		2
1048571	51	2.0	2	...	1		2
1048572	55	2.0	2	...	2		2
1048573	28	2.0	2	...	2		2
1048574	52	2.0	2	...	2		2

	CARDIOVASCULAR	OBESITY	RENAL_CHRONIC	TOBACCO
CLASIFFICATION_FINAL ICU \				
0	2	2	2	2
3 2				
1	2	1	1	2
5 2				
2	2	2	2	2
3 2				
3	2	2	2	2
7 2				
4	2	2	2	2
3 2				

...
...
1048570	2	2	2	2
7 2				
1048571	2	2	2	2
7 2				
1048572	2	2	2	2
7 2				
1048573	2	2	2	2
7 2				
1048574	2	2	2	2

```
7    2
```

	Alive	Age Group
0	2	5.0
1	2	5.0
2	2	4.0
3	2	4.0
4	2	5.0
...
1048570	1	3.0
1048571	1	4.0
1048572	1	4.0
1048573	1	2.0
1048574	1	4.0

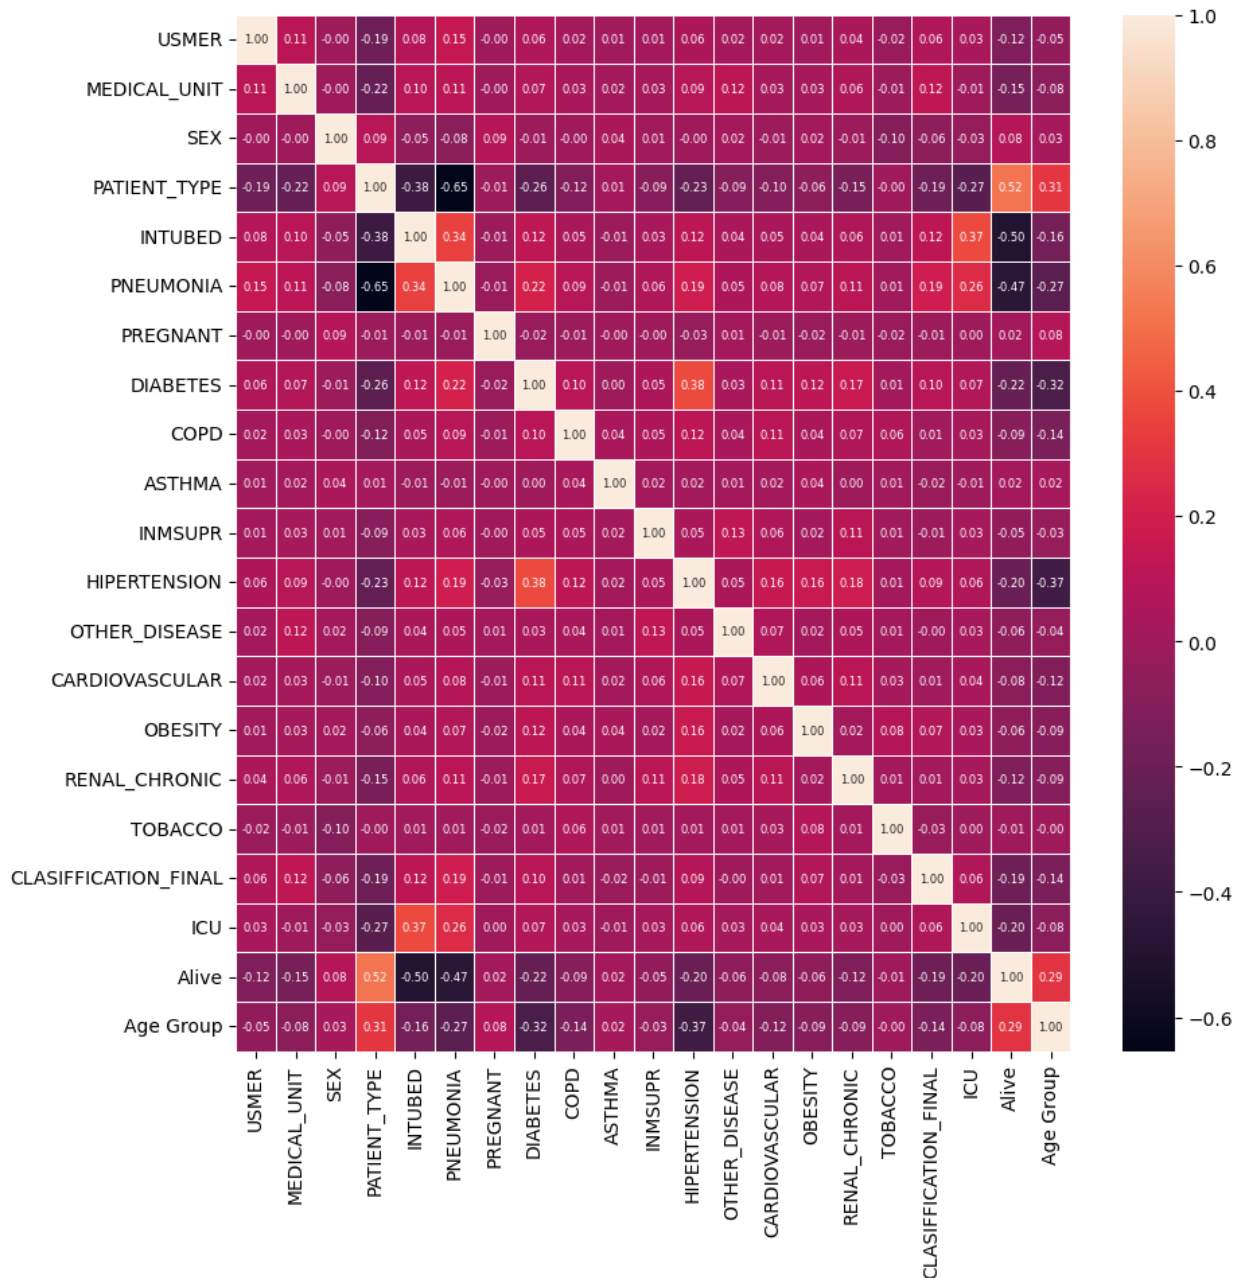
```
[1019666 rows x 23 columns]
```

```
refined_data.drop(columns=['DATE_DIED', 'AGE'], axis=1, inplace=True)
```

Observations

- With respect to 'Alive' we have 'PNEUMONIA', 'INTUBED', 'PATIENT_TYPE', 'Age Group' are most prominent features.
- But as the dataset is unbalanced corr might not help much in feature selection.

```
plt.figure(figsize=(10, 10))
sns.heatmap(refined_data.corr(method='spearman'), annot=True,
fmt=".2f", linewidths=0.5, annot_kws={"size": 6})
plt.show()
```



```
discrete_features = refined_data.columns
```

Observation

- We could see it is difficult to make prediction for 'Alive' as value count is almost similar in ratio with all classes.
- 'PNEUMONIA', 'INTUBED', 'Age Group' shows a major differentiation.

```
# Define the number of features and set a dynamic height
num_features = len(discrete_features)
fig, ax = plt.subplots(num_features, 2, figsize=(16, num_features * 2))
```



```

2.5)) # Increase height dynamically

for i, feature in enumerate(discrete_features):
    # Plot countplot without hue
    sns.countplot(ax=ax[i, 0], x=feature, data=refined_data)
    ax[i, 0].set_title(f"Distribution of {feature}", fontsize=12)
    ax[i, 0].set_xlabel("")
    ax[i, 0].set_ylabel("Count", fontsize=10)



















    # Plot countplot with 'Alive' hue
    sns.countplot(ax=ax[i, 1], x=feature, hue='Alive',
data=refined_data)
    ax[i, 1].set_title(f"{feature} vs Alive", fontsize=12)
    ax[i, 1].set_xlabel("")
    ax[i, 1].set_ylabel("Count", fontsize=10)

    # Set y-axis ticks dynamically
    max_count = refined_data[feature].value_counts().max()
    tick_interval = max(1, max_count // 10) # Set dynamic interval
for better spacing
    ax[i, 0].set_yticks(range(0, max_count + tick_interval,
tick_interval))

    # Set y-axis tick spacing dynamically for the second plot (based
on hue)
    max_count_2 = refined_data.groupby(feature)
['Alive'].value_counts().max()
    tick_interval_2 = max(1, max_count_2 // 10)
    ax[i, 1].set_yticks(range(0, max_count_2 + tick_interval_2,
tick_interval_2))

plt.subplots_adjust(hspace=0.5, wspace=0.3) # Increase spacing
between subplots
fig.tight_layout(pad=2)
plt.show()

```



Observation

- We could see the top features

Age Group : 43822.594537976875
MEDICAL_UNIT : 37796.72489121804
PATIENT_TYPE : 34859.798318027286
CLASIFFICATION_FINAL : 26083.104291404423
PNEUMONIA : 14077.887325773387
INTUBED : 4130.213961265305
HIPERTENSION : 2994.5650815717036
DIABETES : 2649.2413529553733
USMER : 1973.9843896470545
SEX : 1085.1027848188737

```
x = refined_data.drop(['Alive'], axis=1).astype(int)
y = refined_data['Alive'].astype(int)
y = y.apply(lambda v: v-1)
best_features = SelectKBest(chi2, k=10)
features_ranking = best_features.fit(x, y)
ranking_dictionary = {}
for i in range(len(features_ranking.scores_)):
    ranking_dictionary[x.columns[i]] = features_ranking.scores_[i]

asc_sort = sorted(ranking_dictionary.items(), key = lambda
x:x[1],reverse=True)

for feature, score in asc_sort:
    print(feature, ': ', score)
```

Age Group : 43822.594537976875
MEDICAL_UNIT : 37796.72489121804
PATIENT_TYPE : 34859.798318027286
CLASIFFICATION_FINAL : 26083.104291404423
PNEUMONIA : 14077.887325773387
INTUBED : 4130.213961265305
HIPERTENSION : 2994.5650815717036
DIABETES : 2649.2413529553733
USMER : 1973.9843896470545
SEX : 1085.1027848188737
ICU : 338.67464097467257
OBESITY : 219.7056593954017
RENAL_CHRONIC : 129.00281808135935
CARDIOVASCULAR : 57.432984168373984
COPD : 56.831732330119344
OTHER_DISEASE : 43.61760073301738
INMSUPR : 16.639369764568322
ASTHMA : 4.673685536567792
PREGNANT : 1.7195691419068941
TOBACCO : 1.2481071577055332

Observation and Insights

- Missing data was easily interpretable
- With increasing 'Age Group' death chances increases.
- Patient with 'PNEUMONIA', 'INTUBED' are having higher chances of death.
- Patients with 'Covid_classification' as '3' are highly death prone than any other classification. That means Covid with level 3 is the most Severe and damaging.
- Overall people who had pre-diagnosed diseases have higher chances of death.
- Patients who have been hospital had severe covid condition leading to more deaths.

Model Building and Hyper Tuning

- Using Logistic Regression, Decision Classifier and Random Forest models for training.
- Using XGBClassifier to apply boosting for unbalanced dataset.
- as it has too many features and values KNN, and SVC are too slow and complex to fit but could use in future with selected features.
- Using SMOTE technique for resampling the class distribution.
- Using F1 score, Recall as primary metrics as the dataset is unbalanced

```
# Stratified train-test split (80-20 split)
X_train, X_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, stratify=y, random_state=123)

# Check class distribution in train & test sets
print("Train: ", y_train.value_counts(normalize=True))
print("Test: ", y_test.value_counts(normalize=True))

Train:  Alive
0      0.928104
1      0.071896
Name: proportion, dtype: float64
Test:  Alive
0      0.928104
1      0.071896
Name: proportion, dtype: float64

lr_model = LogisticRegression(max_iter=1000,
random_state=123, verbose=1)
dt_model = DecisionTreeClassifier(max_depth=25, max_features=10,
random_state=123)
rf_model = RandomForestClassifier(n_estimators=150, max_depth=25,
random_state=123, verbose=1)
xgb_model = XGBClassifier(n_estimators=150, random_state=123)

X_train= X_train.astype(int)
X_test= X_test.astype(int)
y_test = y_test.astype(int)
y_train= y_train.astype(int)

lr_model.fit(X_train,y_train)
lr_y_pred= lr_model.predict(X_test)
```

```
print('Logistic Regression Complete')
```

```
xgb_model.fit(X_train,y_train,verbose=1)
```

```
xgb_y_pred= xgb_model.predict(X_test)
```

```
print('XGB Complete')
```

```
dt_model.fit(X_train,y_train)
```

```
dt_y_pred= dt_model.predict(X_test)
```

```
print('Decision Tree Regression Complete')
```

```
rf_model.fit(X_train,y_train)
```

```
rf_y_pred= rf_model.predict(X_test)
```

```
print('Random Forrest Complete')
```

Logistic Regression Complete

XGB Complete

Decision Tree Regression Complete

[Parallel(n_jobs=1)]: Done 49 tasks | elapsed: 50.3s

[Parallel(n_jobs=1)]: Done 49 tasks | elapsed: 1.9s

Random Forrest Complete

```
print('lr_model: ',roc_auc_score(y_test,
lr_model.predict_proba(X_test)[: , 1]),end='\n\n')
```

```
print('dt_model: ',roc_auc_score(y_test,
dt_model.predict_proba(X_test)[: , 1]),end='\n\n')
```

```
print('rf_model: ',roc_auc_score(y_test,
rf_model.predict_proba(X_test)[: , 1]),end='\n\n')
```

```
print('xgb_model: ',roc_auc_score(y_test,
xgb_model.predict_proba(X_test)[: , 1]),end='\n\n')
```

```
print('lr_model: ',classification_report(y_test,lr_y_pred),end='\n\n')
```

```
print('dt_model: ',classification_report(y_test,dt_y_pred),end='\n\n')
```

```
print('rf_model: ',classification_report(y_test,rf_y_pred),end='\n\n')
```

```
print('xgb_model: ',classification_report(y_test,xgb_y_pred),end='\n\n')
```

lr_model: 0.9611906536124379

dt_model: 0.8842313179781945

[Parallel(n_jobs=1)]: Done 49 tasks | elapsed: 2.0s

rf_model: 0.9539316447546056

xgb_model: 0.9676899147520295

lr_model:		precision	recall	f1-score	support
-----------	--	-----------	--------	----------	---------

0	0.96	0.98	0.97	189272	
---	------	------	------	--------	--

1	0.70	0.50	0.58	14662
accuracy			0.95	203934
macro avg	0.83	0.74	0.78	203934
weighted avg	0.94	0.95	0.94	203934
dt_model:				
	precision		recall	f1-score
	support			
0	0.96	0.98	0.97	189272
1	0.68	0.49	0.57	14662
accuracy			0.95	203934
macro avg	0.82	0.74	0.77	203934
weighted avg	0.94	0.95	0.94	203934
rf_model:				
	precision		recall	f1-score
	support			
0	0.96	0.98	0.97	189272
1	0.69	0.52	0.59	14662
accuracy			0.95	203934
macro avg	0.83	0.75	0.78	203934
weighted avg	0.94	0.95	0.95	203934
xgb_model:				
	precision		recall	f1-score
	support			
0	0.96	0.99	0.97	189272
1	0.73	0.52	0.61	14662
accuracy			0.95	203934
macro avg	0.85	0.75	0.79	203934
weighted avg	0.95	0.95	0.95	203934

```

smote = SMOTE(sampling_strategy=0.5, random_state=123)
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)

lr_model.fit(X_resampled,y_resampled)
lr_y_pred= lr_model.predict(X_test)
print('Logistic Regression Complete')

xgb_model.fit(X_resampled,y_resampled,verbose=1)
xgb_y_pred= xgb_model.predict(X_test)
print('XGB Complete')

dt_model.fit(X_resampled,y_resampled)
dt_y_pred= dt_model.predict(X_test)
print('Decision Tree Regression Complete')

```

```

rf_model.fit(X_resampled,y_resampled)
rf_y_pred= rf_model.predict(X_test)
print('Random Forrest Complete')

print('lr_model: ',roc_auc_score(y_test,
lr_model.predict_proba(X_test)[: , 1]),end='\n\n')
print('dt_model: ',roc_auc_score(y_test,
dt_model.predict_proba(X_test)[: , 1]),end='\n\n')
print('rf_model: ',roc_auc_score(y_test,
rf_model.predict_proba(X_test)[: , 1]),end='\n\n')
print('xgb_model: ',roc_auc_score(y_test,
xgb_model.predict_proba(X_test)[: , 1]),end='\n\n')

print('lr_model: ',classification_report(y_test,lr_y_pred),end='\n\n')
print('dt_model: ',classification_report(y_test,dt_y_pred),end='\n\n')
print('rf_model: ',classification_report(y_test,rf_y_pred),end='\n\n')
print('xgb_model: ',classification_report(y_test,xgb_y_pred),end='\n\n')

```

Logistic Regression Complete

XGB Complete

Decision Tree Regression Complete

```

[Parallel(n_jobs=1)]: Done 49 tasks      | elapsed: 1.3min
[Parallel(n_jobs=1)]: Done 49 tasks      | elapsed: 1.9s

```

Random Forrest Complete

lr_model: 0.9611741113618207

dt_model: 0.8876410984628946

```

[Parallel(n_jobs=1)]: Done 49 tasks      | elapsed: 2.2s

```

rf_model: 0.950391577537917

xgb_model: 0.9672017062047686

lr_model:		precision	recall	f1-score	support
	0	0.99	0.92	0.95	189272
	1	0.45	0.86	0.59	14662
	accuracy			0.92	203934
	macro avg	0.72	0.89	0.77	203934
	weighted avg	0.95	0.92	0.93	203934

dt_model:		precision	recall	f1-score	support
-----------	--	-----------	--------	----------	---------

	0	0.98	0.93	0.95	189272
	1	0.46	0.82	0.59	14662
accuracy				0.92	203934
macro avg		0.72	0.87	0.77	203934
weighted avg		0.95	0.92	0.93	203934
rf_model:					
		precision	recall	f1-score	support
	0	0.99	0.93	0.95	189272
	1	0.46	0.83	0.60	14662
accuracy				0.92	203934
macro avg		0.73	0.88	0.78	203934
weighted avg		0.95	0.92	0.93	203934
xgb_model:					
		precision	recall	f1-score	support
	0	0.99	0.92	0.95	189272
	1	0.45	0.90	0.60	14662
accuracy				0.91	203934
macro avg		0.72	0.91	0.78	203934
weighted avg		0.95	0.91	0.93	203934

```

class_weight= {0:9.2,1:0.8}
lr_model = LogisticRegression(max_iter=1000,
random_state=123,verbose=1,class_weight=class_weight)
dt_model = DecisionTreeClassifier(max_depth=25, max_features=10,
random_state=123,class_weight=class_weight)
rf_model = RandomForestClassifier(n_estimators=150, max_depth=25,
random_state=123,verbose=1,class_weight=class_weight)
xgb_model = XGBClassifier(n_estimators=150, random_state=123)

lr_model.fit(X_train,y_train)
lr_y_pred= lr_model.predict(X_test)
print('Logistic Regression Complete')

dt_model.fit(X_train,y_train)
dt_y_pred= dt_model.predict(X_test)
print('Decision Tree Regression Complete')
rf_model.fit(X_train,y_train)
rf_y_pred= rf_model.predict(X_test)
print('Random Forrest Complete')

print('lr_model: ',roc_auc_score(y_test,
lr_model.predict_proba(X_test)[:,: 1]),end='\n\n')

```



```

print('dt_model: ',roc_auc_score(y_test,
dt_model.predict_proba(X_test)[: , 1]),end='\n\n')
print('rf_model: ',roc_auc_score(y_test,
rf_model.predict_proba(X_test)[: , 1]),end='\n\n')
# print('xgb_model: ',roc_auc_score(y_test,
xgb_model.predict_proba(X_test)[: , 1]),end='\n\n')

```

Logistic Regression Complete
Decision Tree Regression Complete

```

[Parallel(n_jobs=1)]: Done 49 tasks      | elapsed: 51.9s
[Parallel(n_jobs=1)]: Done 49 tasks      | elapsed: 2.2s

```

Random Forrest Complete
lr_model: 0.9607487445568135

dt_model: 0.886563029397784

```

[Parallel(n_jobs=1)]: Done 49 tasks      | elapsed: 3.0s
rf_model: 0.9522905035171297

```

```

print('lr_model: ',classification_report(y_test,lr_y_pred),end='\n\n')
print('dt_model: ',classification_report(y_test,dt_y_pred),end='\n\n')
print('rf_model: ',classification_report(y_test,rf_y_pred),end='\n\n')
print('xgb_model: ',classification_report(y_test,xgb_y_pred),end='\n\n')

```

lr_model:		precision	recall	f1-score	support
	0	0.93	1.00	0.97	189272
	1	0.93	0.08	0.15	14662
accuracy			0.93		203934
macro avg	0.93	0.54	0.56		203934
weighted avg	0.93	0.93	0.91		203934

dt_model:		precision	recall	f1-score	support
	0	0.94	0.99	0.97	189272
	1	0.70	0.24	0.36	14662
accuracy			0.94		203934
macro avg	0.82	0.62	0.66		203934
weighted avg	0.93	0.94	0.92		203934

rf_model:		precision	recall	f1-score	support
	0	0.95	0.99	0.97	189272
	1	0.76	0.27	0.40	14662
accuracy			0.94	203934	
macro avg		0.85	0.63	0.68	203934
weighted avg		0.93	0.94	0.93	203934
xgb_model:		precision	recall	f1-score	support
	0	0.99	0.92	0.95	189272
	1	0.45	0.90	0.60	14662
accuracy			0.91	203934	
macro avg		0.72	0.91	0.78	203934
weighted avg		0.95	0.91	0.93	203934

Best Model To Choose

- XGBClassifier with Smote is the best model to choose due to high Accuracy 91 and recall 91, as we could go with False positive in this case rather than taking risk.

Future Suggestion/Scope

- Should work on a better trade off balance for minority class in this case.
- Should Focus on Hyperparameter tuning for better results.
- More classification could be done such as predicting the level of covid through other features.