

# Unsupervised Learning Project: Spotify Songs Clustering

## Objective

The main objective of this clustering algorithm is to be later on used for similar song recommendation. I will be using Dimensionality reduction with clustering algorithms.

## Import Packages

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import nltk
from wordcloud import WordCloud
from collections import Counter
from textblob import TextBlob
from sklearn.preprocessing import StandardScaler, OneHotEncoder,
Normalizer, MinMaxScaler
from sklearn.manifold import TSNE, MDS
from umap.umap_ import UMAP
import os
from sentence_transformers import SentenceTransformer
from sklearn.decomposition import TruncatedSVD, PCA, KernelPCA
from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
from mpl_toolkits.mplot3d import Axes3D
import hdbscan
from sklearn.metrics import silhouette_score
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn
warnings.filterwarnings('ignore')
```

## Dataset

### [Dataset Link](#)

This dataset contains a collection of the most popular songs on Spotify, along with various attributes that can be used for music analysis and recommendation systems. It includes audio features, lyrical details, and general metadata about each track, making it an excellent resource for machine learning, data science, and music analytics projects.

Each song in the dataset includes the following features:

#### 🔊 Audio Features (Extracted from Spotify API):

- Danceability – How suitable a track is for dancing (0.0 – 1.0).
- Energy – Intensity and activity level of a song (0.0 – 1.0).
- Loudness – Overall loudness in decibels (dB).
- Speechiness – Presence of spoken words in the track (0.0 – 1.0).
- Acousticness – Probability that a track is acoustic (0.0 – 1.0).
- Instrumentalness – Predicts if a track is instrumental (0.0 – 1.0).
- Liveness – Probability of a live audience (0.0 – 1.0).
- Valence – Musical positivity or happiness (0.0 – 1.0).
- Tempo – Beats per minute (BPM) of the track.
- Key & Mode – Musical key and mode (major/minor).

#### 📄 Lyrics-Based Features:

- Lyrics Text – Full lyrics of the song (if available).

#### 🎵 General Song Information:

- Track Name – Name of the song.
- Artist(s) – Performing artist(s).
- Album Name – Album the track belongs to.
- Release Year – Year when the song was released.
- Genre – Song's primary genre classification.
- Popularity Score – Spotify popularity metric (0 – 100).

#### Already Preprocessed Columns:

- tags: lyrics stopwords and punctuations has been removed.
- tags\_tokenized: tags has been split into list
- doc\_vector: Word2Vec embedding has been used for getting tokenized words embeddings. size: (300,)
- combined\_vector: 11 Numerical featured has been combined with doc\_vector giving size of (311,)
- cluster : A column has been added by the dataset provider based on his Kmeans clustering. (That I will not be using)
- image\_url: url of the spotify song image
- spotify\_url: track spotify url

## EDA

EDA will be consisting of:

- Column names and unique values with each values
- Pair plot to have a better understanding how each column define different clusters
- Heatmap of numerical columns for understanding impact of each column on one another.
- Looking for NA values and replacing them.

- Finding Important features and reducing dimensions for computational efficiency and finding higher dimensions feature relations.
- Getting Insights from different graphs.
- Has used One Hot Encoding, normalization, standard scaling like techniques.
- Used SentenceTransformer for getting semantic meaning of the lyrics or tags.

## Loading Dataset

```
data = pd.read_csv("songs.csv", index_col= 0)
data.head(3)
```

	track_id	track_name	track_artist
0	6oJ6le65B3SEqPwMRNXWjY	higher love	Kygo
1	3yNZ5r3LKfdmjoS3gkhUCT	bad guy (with justin bieber)	Billieeilish
2	0qc4QlcCxVTGyShurEv1UU	post malone (feat. rani)	Samfeldt

	track_popularity	track_album_release_date	playlist_genre
0	0.500000	2019-06-28	Pop
1	0.318182	2019-07-11	Pop
2	0.318182	2019-05-24	Pop

	energy	key	loudness	...	track_artist_merged
0	0.667346	0.727273	0.680129	...	kygo
1	0.425904	0.000000	0.504094	...	billie eilish
2	0.628716	0.636364	0.821136	...	sam feldt

	lyrics	artist_name
0	['bring', 'me', 'higher', 'love,', 'love', '...', ...]	Kygo
1	['yeah,', 'yeah', '...', 'oh,', 'ah', '...', 'white'...]	Billie Eilish
2	['one', 'more', 'drink,', 'got', 'one', 'more'...]	Sam Feldt

	tags
0	bring higher love love bring higher love think...
1	yeah yeah oh ah white shirt red bloody nose sl...
2	one drink got one bacardi one dance afterparty...

	tags_tokenized
0	['bring', 'higher', 'love', 'love', 'bring', '...', ...]
1	['yeah', 'yeah', 'oh', 'ah', 'white', 'shirt', '...', ...]
2	['one', 'drink', 'got', 'one', 'bacardi', 'one...', ...]

	doc_vector
0	[-0.1148182 0.27755967 0.27891365 0.143460...]

```

1 [-2.95320839e-01 -2.59309914e-03  3.84592146e-...
2 [-0.1616459      0.21872164   0.3755187      0.116120...

                                combined_vector cluster \
0 [-0.1148182      0.27755967   0.27891365   0.143460...      3
1 [-2.95320839e-01 -2.59309914e-03  3.84592146e-...      1
2 [-0.1616459      0.21872164   0.37551871   0.116120...      3

                                image_url \
0 https://i.scdn.co/image/ab67616d0000b2737c8977...
1 https://i.scdn.co/image/ab67616d0000b273a69b8b...
2 https://i.scdn.co/image/ab67616d0000b27354de16...

                                spotify_url
0 https://open.spotify.com/track/6oJ6le65B3SEqPw...
1 https://open.spotify.com/track/3yNZ5r3LKfdmjoS...
2 https://open.spotify.com/track/0qc4QlcCxVTGySh...

[3 rows x 27 columns]

data.drop(columns=['image_url','spotify_url'],axis=1, inplace=True)
data.columns

Index(['track_id', 'track_name', 'track_artist', 'track_popularity',
      'track_album_release_date', 'playlist_genre', 'danceability',
      'energy',
      'key', 'loudness', 'mode', 'speechiness', 'acousticness',
      'liveness',
      'valence', 'tempo', 'duration_ms', 'track_artist_merged',
      'lyrics',
      'artist_name', 'tags', 'tags_tokenized', 'doc_vector',
      'combined_vector', 'cluster'],
      dtype='object')

```

## Observation

- Dropped unnecessary features which will not be used,'image\_url','spotify\_url' as they will not be contributing anything.

```

data.info()

<class 'pandas.core.frame.DataFrame'>
Index: 900 entries, 0 to 899
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   track_id                             900 non-null    object
1   track_name                           900 non-null    object
2   track_artist                         900 non-null    object
3   track_popularity                     900 non-null    float64
4   track_album_release_date             900 non-null    object

```

5	playlist_genre	900	non-null	object
6	danceability	900	non-null	float64
7	energy	900	non-null	float64
8	key	900	non-null	float64
9	loudness	900	non-null	float64
10	mode	900	non-null	float64
11	speechiness	900	non-null	float64
12	acousticness	900	non-null	float64
13	liveness	900	non-null	float64
14	valence	900	non-null	float64
15	tempo	900	non-null	float64
16	duration_ms	900	non-null	int64
17	track_artist_merged	900	non-null	object
18	lyrics	900	non-null	object
19	artist_name	900	non-null	object
20	tags	900	non-null	object
21	tags_tokenized	900	non-null	object
22	doc_vector	900	non-null	object
23	combined_vector	900	non-null	object
24	cluster	900	non-null	int64

dtypes: float64(11), int64(2), object(12)  
memory usage: 182.8+ KB

### Obsevation

- No null values has been found
- doc\_vector and combined\_vector is in string format we have to convert them into np.array.
- we have 11 Numerically Scaled features, 1 integer feature duration

```
for column in data.columns:
    print(column, len(data[column].unique()))
```

track\_id 900  
track\_name 900  
track\_artist 468  
track\_popularity 23  
track\_album\_release\_date 467  
playlist\_genre 6  
danceability 424  
energy 474  
key 12  
loudness 844  
mode 2  
speechiness 577  
acousticness 664  
liveness 462  
valence 550  
tempo 876  
duration\_ms 887

```
track_artist_merged 468
lyrics 899
artist_name 468
tags 900
tags_tokenized 900
doc_vector 900
combined_vector 900
cluster 5
```

## Observation

- There are 468 different artist, we could also see artist\_name, track\_artist\_merged, and track\_artist all have same feature with different preprocessing.
- we have 2 modes
- 6 genre
- 12 keys

```
import ast
def convert_to_numpy_array(array_str):
    # Remove the brackets and split by space, then convert to float
    array_str = array_str.strip('[]') # Remove the square brackets
    return np.array([float(x) for x in array_str.split()]) # Convert to float and create a NumPy array

# Convert the specified columns to NumPy arrays
data['doc_vector'] = data['doc_vector'].apply(convert_to_numpy_array)
data['combined_vector'] =
data['combined_vector'].apply(convert_to_numpy_array)

print('combined_vector', 'doc_vector')
for i in range(10):
    print(len(data['combined_vector'][i]), len(data['doc_vector'][i]))

combined_vector doc_vector
311 300
311 300
311 300
311 300
311 300
311 300
311 300
311 300
311 300
311 300
```

## Observation

- We get to know combined vector is a mix of 11 floating scaled values and doc\_vector as 11 values are only added to it and we could also verify that by comparing starting 300 values with doc\_vector and last 11 with the floating scaled numerical column data

```

# Selecting only numerical columns
numerical_cols = list(data.select_dtypes(include=['number']).columns)
categorical_cols = list(data.columns.difference(numerical_cols))
categorical_cols, numerical_cols

(['artist_name',
 'combined_vector',
 'doc_vector',
 'lyrics',
 'playlist_genre',
 'tags',
 'tags_tokenized',
 'track_album_release_date',
 'track_artist',
 'track_artist_merged',
 'track_id',
 'track_name'],
 ['track_popularity',
 'danceability',
 'energy',
 'key',
 'loudness',
 'mode',
 'speechiness',
 'acousticness',
 'liveness',
 'valence',
 'tempo',
 'duration_ms',
 'cluster'])

text = "".join(data['track_name'])
wordcloud = WordCloud(width = 600, height=300, background_color =
'white').generate(text)
plt.figure(figsize = (6,3))
plt.imshow(wordcloud,interpolation = 'bilinear')
plt.axis("off")
plt.title("Most Common Words in Track Names")
plt.show()

```

- Feat, Love, Drake are most used words and we could also notice the other most common words for the songs.

```
# Count of songs released per year
data['track_album_release_date'] =
pd.to_datetime(data['track_album_release_date'], errors='coerce')
data['year'] = data['track_album_release_date'].dt.year
yearly_songs = data.groupby('year')['track_id'].count()
# Define bin size (change to 5, 10, etc., based on dataset size)
bin_size = 5

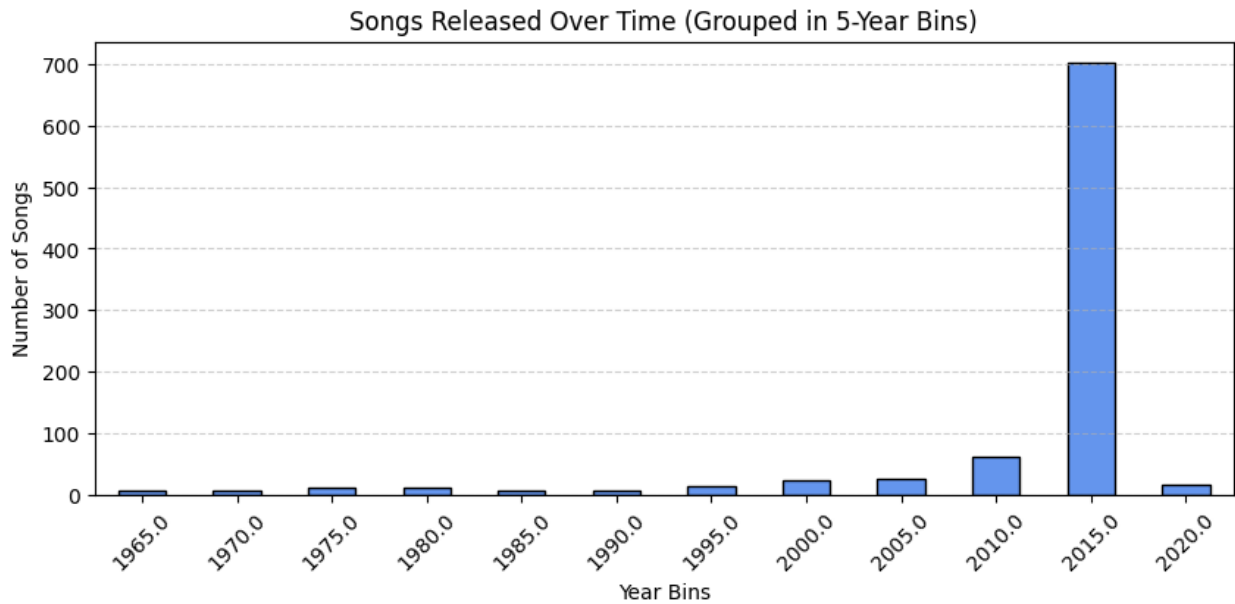
# Create bins and labels
data['year_bin'] = (data['year'] // bin_size) * bin_size # Groups
into bins like 1990-1995, 1995-2000, etc.

# Count songs per bin
yearly_songs_binned = data.groupby('year_bin')['track_id'].count()

# Plot
plt.figure(figsize=(10, 4))
yearly_songs_binned.plot(kind="bar", color="cornflowerblue",
edgecolor="black")

plt.title(f"Songs Released Over Time (Grouped in {bin_size}-Year
Bins)")
plt.xlabel("Year Bins")
plt.ylabel("Number of Songs")
plt.xticks(rotation=45) # Rotate labels for better readability
plt.grid(axis="y", linestyle="--", alpha=0.6)
plt.show()
```

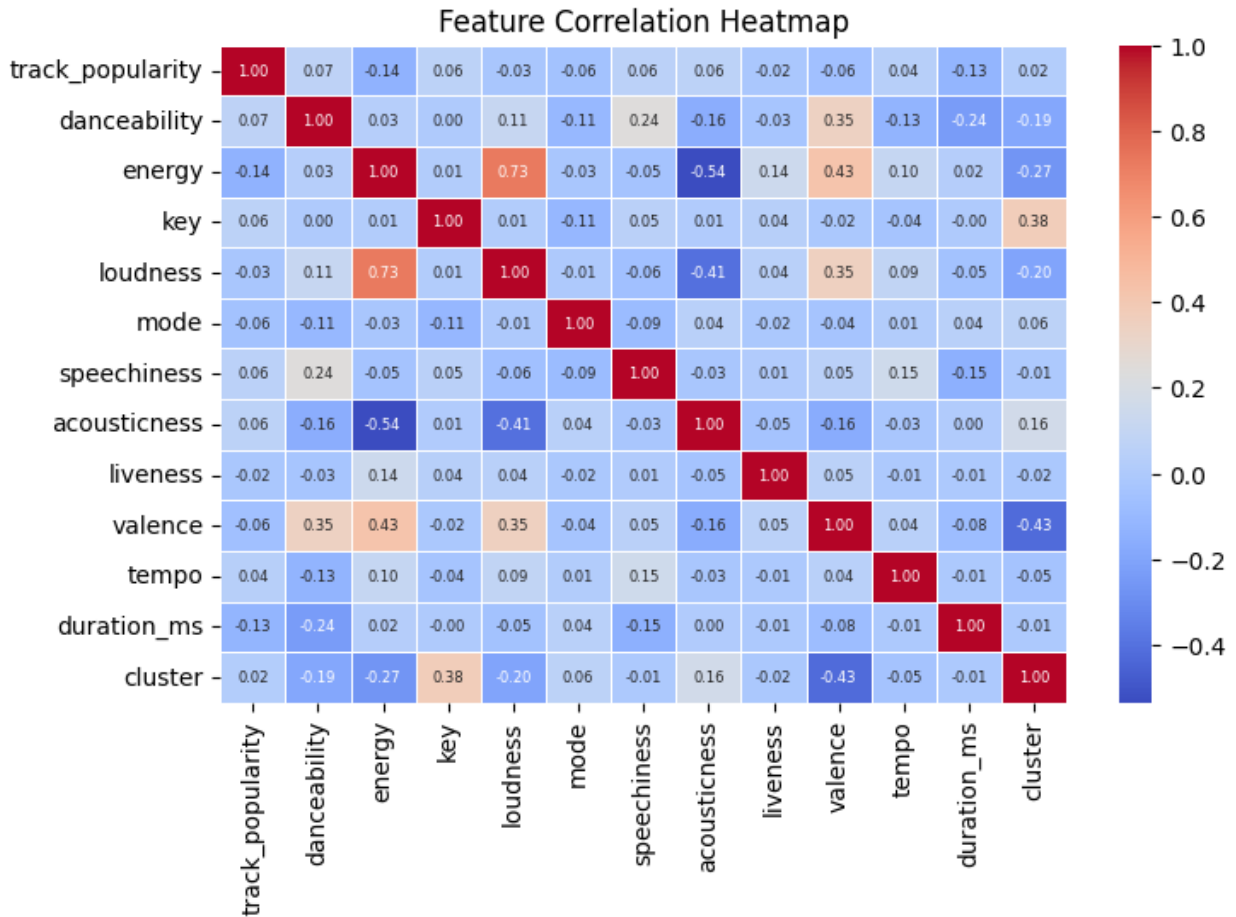




#### Observation

- Most songs has been released in **2015-2020** spam i.e **700**
- We could also see increasing trend of songs and emerging artist through it.

```
plt.figure(figsize=(8,5))
sns.heatmap(data[numerical_cols].corr(), annot=True, cmap="coolwarm",
fmt=".2f",linewidths=0.5, annot_kws={"size": 6})
plt.title("Feature Correlation Heatmap")
plt.show()
```

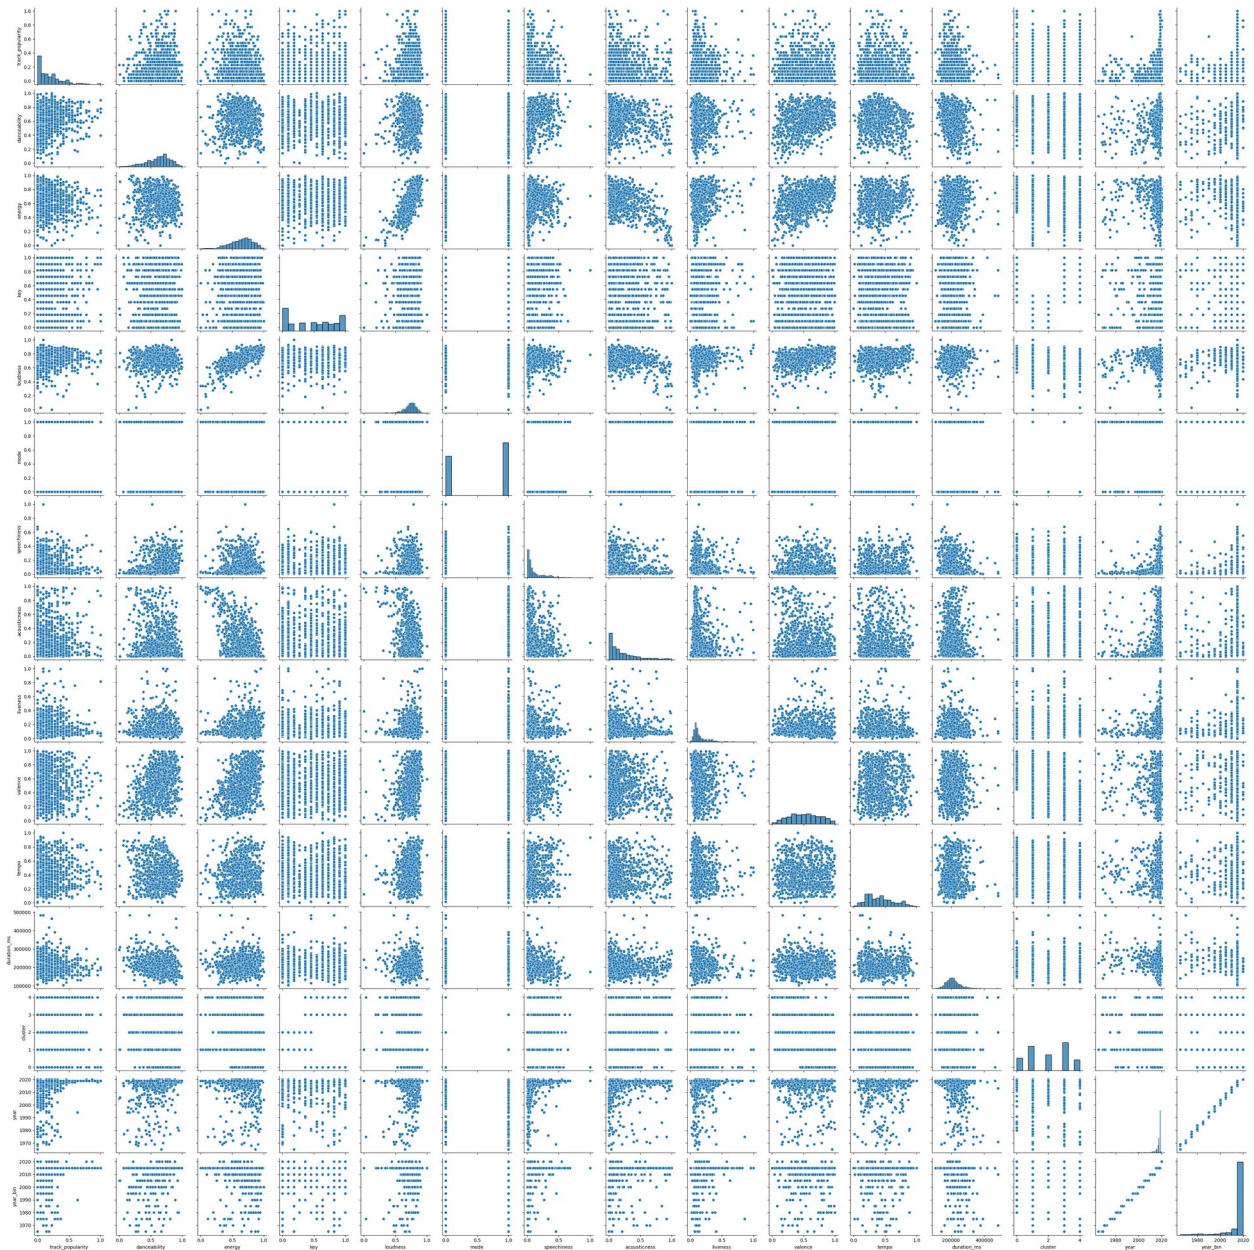


#### Observation

- Loudness is highly related positively with Energy, and negatively with acousticness
- Danceability is highly related with valence
- Key and valence plays an important role in clustering for the previous model.

```
sns.pairplot(data.drop(columns=['doc_vector', 'combined_vector'], axis=1))
```

```
<seaborn.axisgrid.PairGrid at 0x2989103acc0>
```



## Observation

- With 2D it won't be easy to cluster with have to go to higher dimentionts.

```
## Seperating Combined vector for analysis.
embedding_matrix = np.vstack(data["combined_vector"].values)
print(embedding_matrix.shape) # (n, 311)
embedding_matrix

(900, 311)

array([[ -0.1148182 ,  0.27755967,  0.27891365, ...,  0.08600444,
         0.39136969,  0.29060535],
       [ -0.29532084, -0.0025931 ,  0.38459215, ...,  0.10293029,
```

```

    0.68763418, 0.50837382],
    [-0.1616459, 0.21872164, 0.37551871, ..., 0.0902359,
    0.65650494, 0.31443855],
    ...,
    [-0.12616959, 0.19998066, 0.29130945, ..., 0.13255051,
    0.35702018, 0.31791831],
    [-0.12194037, 0.32961181, 0.34245756, ..., 0.21083254,
    0.43752684, 0.30454539],
    [-0.08236136, 0.28211612, 0.30290502, ..., 0.17909658,
    0.45899528, 0.61985913]])

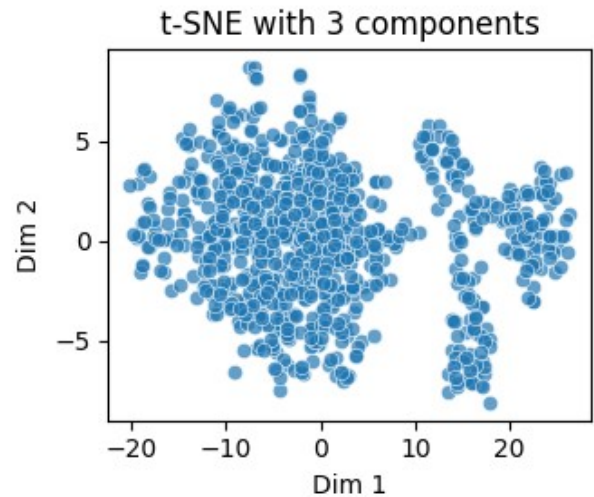
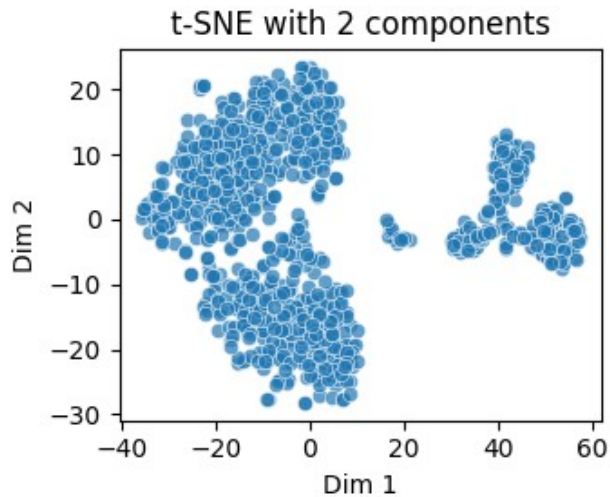
def plot_tsne(X_data):
    fig, axes = plt.subplots(1, 2, figsize=(7, 3))
    axes = axes.ravel() # Flatten the array of axes for easy indexing

    for i, n_comp in enumerate(range(2, 4)):
        # Applying t-SNE
        tsne = TSNE(n_components=n_comp, random_state=42)
        X_tsne = tsne.fit_transform(X_data)

        # Plot only the first two dimensions
        ax = axes[i]
        sns.scatterplot(
            x=X_tsne[:, 0],
            y=X_tsne[:, 1],
            alpha=0.7,
            ax=ax
        )
        ax.set_title(f"t-SNE with {n_comp} components")
        ax.set_xlabel("Dim 1")
        ax.set_ylabel("Dim 2")

    plt.tight_layout()
    plt.show()
plot_tsne(X_data=embedding_matrix)

```

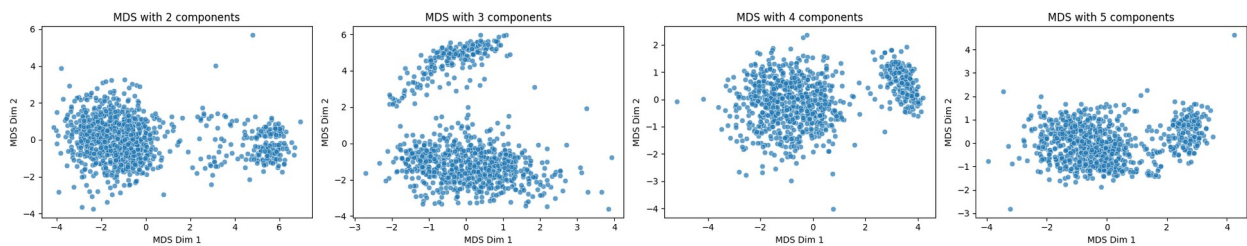


```
def plot_msd(X_data):
    fig, axes = plt.subplots(1, 4, figsize=(20, 4))
    axes = axes.ravel() # Flatten the array of axes for easy indexing

    for i, n_comp in enumerate(range(2, 6)):
        # Applying t-SNE
        mds = MDS(n_components=n_comp, random_state=42)
        X_mds = mds.fit_transform(embedding_matrix)

        ax = axes[i]
        sns.scatterplot(
            x=X_mds[:, 0], # First dimension
            y=X_mds[:, 1], # Second dimension
            palette='viridis',
            alpha=0.7,
            ax=ax
        )
        ax.set_title(f"MDS with {n_comp} components")
        ax.set_xlabel("MDS Dim 1")
        ax.set_ylabel("MDS Dim 2")

    plt.tight_layout()
    plt.show()
plot_msd(embedding_matrix)
```





## Observation

- We could have 2-4 clusters in our data based on 2D plot w.r.t. combined\_vector with both dimensionality reduction techniques.
- But we have to keep in mind that MDS and TSNE does not work well with semantic data, they are best for distanced based algorithm.

## Observation And Insights

- Energy, Acousticness, Loudness are inter-related as mellow songs would have more acoustic sounds, less loudness and energy usually
- high loudness results in more energy in songs
- 6 types of genre, 12 types of key or notes are there
- We would need higher dimensions and understanding of song features for similar song recommendation and clustering as 2D won't be much useful.

## Pre-Processing and, Model scoring and defining

- **genre, artist** has been **One Hot Encoded** using `pd.get_dummies`.
- Numerical Data has been separated.
- To get semantic features **SentenceTransformer** is being used with **tags** then **normalised** using **l2**.
- Numerical data is linear so used PCA having 5 component.
- Hot Encoded values has been reduced to 20 for artist and 1 for genre using **TruncatedSVD**.
- **UMAP** is used to reduced features to 30 for getting semantic relation with higher dimensions as there were 300+ tags.

I created 2 Data:

- One with without reducing any features just normalised and scaled.
- 2nd with reduced with different techniques.

```
def plot_kmeans(X_data,name):
    intertias = []
    for k in range(2, 20):
        kmeans = KMeans(n_clusters=k, random_state=42)
        kmeans.fit(X_data)
        intertias.append(kmeans.inertia_)
    plt.figure(figsize=(3,3))
    plt.plot(range(2, 20), intertias)
    plt.xlabel("Number of Clusters (k)")
    plt.ylabel("Intertia")
    plt.title(f"Elbow Method for Optimal k for {name}")
    plt.show()

def DBSCAN_train(X_data,name):
    # Example values (Replace with optimal found from K-Distance plot)
    eps_values = [0.05, 0.1, 0.2, 0.3,0.4,0.5,0.7,0.8] # Adjust based
on dataset
    min_samples_values = [5, 10, 20,50,80,100] # Try values based on
```

```

data density

best_score = -1
best_params = None
for eps in eps_values:
    for min_samples in min_samples_values:
        dbscan = DBSCAN(eps=eps, min_samples=min_samples,
metric='euclidean') # Change metric if needed
        labels = dbscan.fit_predict(X_data)

        # Ignore if only one cluster is formed (or all noise)
        if len(set(labels)-{1}) > 1:
            score = silhouette_score(X_data, labels)
            # print(f"eps={eps}, min_samples={min_samples},
silhouette_score={score:.5f}, clusters={len(set(labels) - {1})}")

            if score > best_score:
                best_score = score
                best_params = (eps, min_samples)
                best_cluster_size= len(set(labels) - {1})
        if best_params:
            print(f"\n{name} Best DBSCAN Params: eps={best_params[0]},
min_samples={best_params[1]}, silhouette_score={best_score:.3f},
Cluster:{best_cluster_size}")

def AgglomerativeClustering_plot(X_data, name):
    scores=[]
    n_comps=[]
    for n in range(2,50):
        agg= AgglomerativeClustering(n_clusters=n)
        cluster_labels = agg.fit_predict(X_data)

        # Ignore Silhouette Score for DBSCAN if many points are marked
as noise (-1)
        if len(set(cluster_labels)) > 1:
            score = silhouette_score(X_data, cluster_labels)
            n_comps.append(n)
            scores.append(score)

    plt.figure(figsize=(3,3))
    plt.plot(n_comps,scores)
    plt.xlabel("Number of Clusters (k)")
    plt.ylabel("Silhouette Score")
    plt.title(f"Scores for Optimal k for {name}")
    plt.show()

#Data Preprocessing
model = SentenceTransformer('all-MiniLM-L6-v2')

df_numerical = data[numerical_cols]

```

```

pca_numerical = PCA(n_components=5, random_state=42)
numerical_reduced = pca_numerical.fit_transform(df_numerical)

# Handling Categorical Columns (One-Hot Encoding 'playlist_genre')
df_genre = pd.get_dummies(data['playlist_genre'], drop_first=True)
df_artist_ohe = pd.get_dummies(data['artist_name'], prefix='artist',
drop_first=True)
artist_svd_onehot = TruncatedSVD(n_components=20, random_state=42)
artist_reduced = artist_svd_onehot.fit_transform(df_artist_ohe)
genre_svd_onehot = TruncatedSVD(n_components=1, random_state=42)
genre_reduced = genre_svd_onehot.fit_transform(df_genre)

# Handling Text Columns using TF-IDF (lyrics, tags)
tag_embeddings = np.vstack(data['tags'].apply(lambda x:
model.encode(str(x))))
normalizer = Normalizer(norm='l2')
embedding_normalized = normalizer.fit_transform(tag_embeddings)
umap_embeddings = UMAP(n_components=30, random_state=42,
metric='cosine')
embeddings_reduced = umap_embeddings.fit_transform(tag_embeddings)

# Convert to DataFrame
df_tags = pd.DataFrame(embedding_normalized, columns=[f'tag_{i}' for i
in range(embedding_normalized.shape[1])])

#Unreduced Data
df_final = pd.concat([df_numerical, df_genre, df_artist_ohe, df_tags],
axis=1)

#Reduced Data
X = np.concatenate([
    numerical_reduced,
    genre_reduced,
    artist_reduced,
    embeddings_reduced
], axis=1)

print(f"Final Shape After Merging: {X.shape}")
df_final.describe()

```

Final Shape After Merging: (900, 56)

	track_popularity	danceability	energy	key
loudness \				
count	900.000000	900.000000	900.000000	900.000000
900.000000				
mean	0.199545	0.623394	0.640775	0.484444



0.726726				
std	0.199525	0.183919	0.181130	0.334353
0.111564				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	0.045455	0.504248	0.527578	0.181818
0.681972				
50%	0.136364	0.652288	0.664127	0.454545
0.743644				
75%	0.318182	0.749346	0.777873	0.818182
0.798564				
max	1.000000	1.000000	1.000000	1.000000
1.000000				

	mode	speechiness	acousticness	liveness
valence ... \				
count 900.000000	900.000000	900.000000	900.000000	900.000000
900.000000 ...				
mean 0.574444	0.115008	0.225910	0.160050	
0.518995 ...				
std 0.494702	0.133004	0.238884	0.139023	
0.237509 ...				
min 0.000000	0.000000	0.000000	0.000000	
0.000000 ...				
25% 0.000000	0.023636	0.043808	0.077727	
0.334210 ...				
50% 1.000000	0.055262	0.140391	0.107691	
0.515887 ...				
75% 1.000000	0.158157	0.326666	0.194965	
0.705882 ...				
max 1.000000	1.000000	1.000000	1.000000	
1.000000 ...				

	tag_374	tag_375	tag_376	tag_377	tag_378
tag_379 \					
count 900.000000	900.000000	900.000000	900.000000	900.000000	900.000000
900.000000					
mean 0.032464	-0.001176	-0.000180	0.083935	-0.004169	
0.008147					
std 0.034599	0.043137	0.034469	0.035275	0.035168	
0.034403					
min -0.073599	-0.125999	-0.128981	-0.057180	-0.111977	-
0.102823					
25% 0.009423	-0.031147	-0.023066	0.060304	-0.027868	-
0.014743					
50% 0.033452	-0.002556	-0.000836	0.083880	-0.004145	
0.006997					
75% 0.057569	0.030092	0.021599	0.109180	0.020649	
0.032362					

max	0.154743	0.136306	0.124025	0.207000	0.116897
0.106140					

	tag_380	tag_381	tag_382	tag_383
count	900.000000	900.000000	900.000000	900.000000
mean	0.038920	0.015560	0.002120	-0.067236
std	0.040891	0.047144	0.041019	0.038570
min	-0.061633	-0.103175	-0.184510	-0.191904
25%	0.009916	-0.019485	-0.025272	-0.094228
50%	0.037997	0.012432	0.002483	-0.070262
75%	0.066414	0.047103	0.031173	-0.043726
max	0.193916	0.158755	0.114898	0.060328

[8 rows x 397 columns]

## Observation

- Non reduced featured has much variance in data, with numerical features scaled from 0-1 and semantic ones are very low.

```
# Standardizing the Data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Using kpca on Non Reduced set with different kernels

kpca = KernelPCA(n_components=50, kernel='poly', gamma=0.3,
random_state=42)
pol_X_kpca = kpca.fit_transform(df_final)
kpca = KernelPCA(n_components=25, kernel='cosine', gamma=0.3,
random_state=42)
cos_X_kpca = kpca.fit_transform(df_final)
kpca = KernelPCA(n_components=50, kernel='rbf', gamma=0.1,
random_state=42)
rbf_X_kpca = kpca.fit_transform(df_final)
```

## Modelling, Training, Plotting and Scoring

- Used Kmeans for spherical cluster in higher dimentions with PCA and Kpca.
- Used DBSCAN for irregular shape and noises like data.
- Used HDBSCAN for high density clusters.
- Used AgglomerativeClustering for Tree based clustering approach based on genre and artist.

```
### Reducing every dataset to 3D using UMAP

def plot_3d(X_data,name):
    # Reduce dimensionality to 3D
    umap_3d = UMAP(n_components=3, random_state=42, metric='cosine')
    X_umap_3d = umap_3d.fit_transform(X_data)
```

```

# Scatter plot in 3D
fig = plt.figure(figsize=(7, 3))
ax = fig.add_subplot(111, projection='3d')

hdb = hdbscan.HDBSCAN(min_cluster_size=2, metric="euclidean")
hdb_labels = hdb.fit_predict(X_umap_3d)

if len(set(hdb_labels)) > 1:
    sil_score = silhouette_score(X_umap_3d, hdb_labels)
    print(name, "Silhouette Score (Ignoring Noise):", sil_score,
"\nclusters:", len(set(hdb_labels)) )

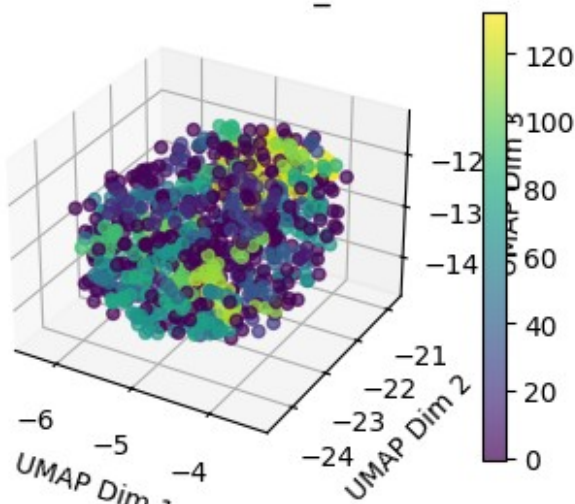
    scatter = ax.scatter(X_umap_3d[:, 0], X_umap_3d[:, 1],
X_umap_3d[:, 2], c=hdb_labels, cmap="viridis", alpha=0.7)
    ax.set_title(f"3D Visualization {name}")
    ax.set_xlabel("UMAP Dim 1")
    ax.set_ylabel("UMAP Dim 2")
    ax.set_zlabel("UMAP Dim 3")
    plt.colorbar(scatter)
    plt.show()

### Different pre-processed data for getting better result and
approach
plot_3d(df_final, 'df_final')
plot_3d(embedding_matrix, 'embedding_matrix')
plot_3d(X_scaled, 'X_scaled')
plot_3d(X, 'X')
plot_3d(pol_X_kpca, 'pol_X_kpca')
plot_3d(cos_X_kpca, 'cos_X_kpca')
plot_3d(rbf_X_kpca, 'rbf_X_kpca')

df_final Silhouette Score (Ignoring Noise): -0.04516269
clusters: 134

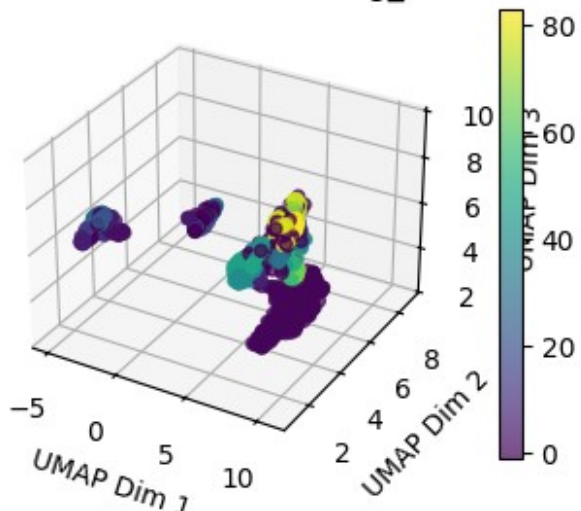
```

3D Visualization df\_final



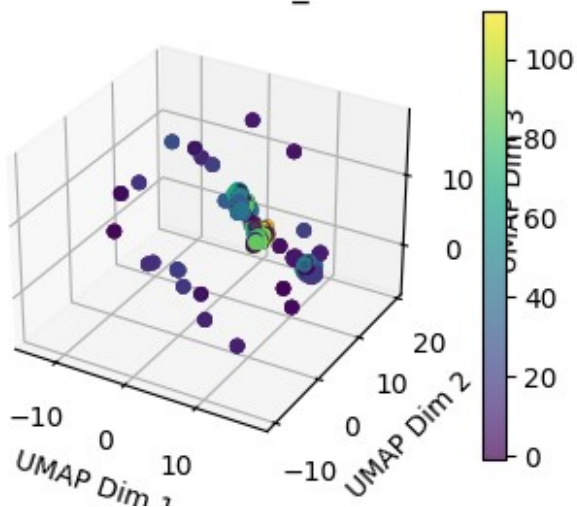
embedding\_matrix Silhouette Score (Ignoring Noise): 0.298826  
clusters: 85

3D Visualization embedding\_matrix



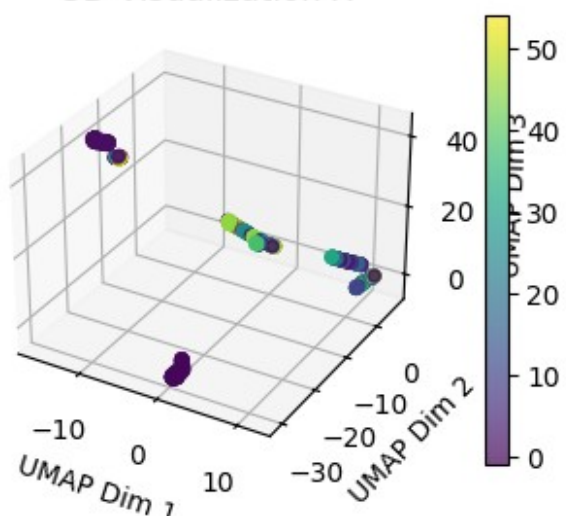
X\_scaled Silhouette Score (Ignoring Noise): 0.4106002  
clusters: 114

3D Visualization X\_scaled



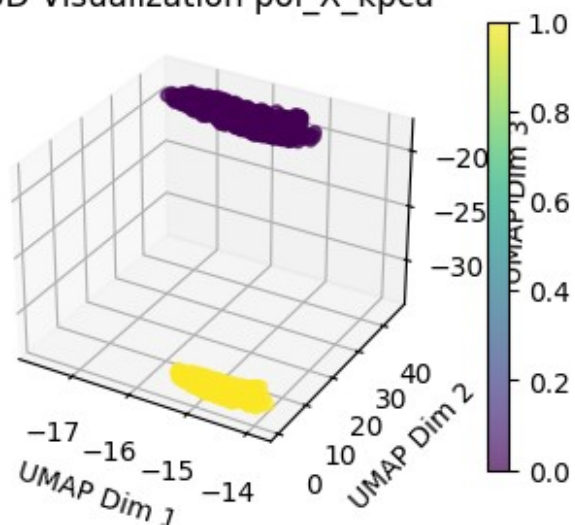
X Silhouette Score (Ignoring Noise): 0.6556491  
clusters: 56

3D Visualization X



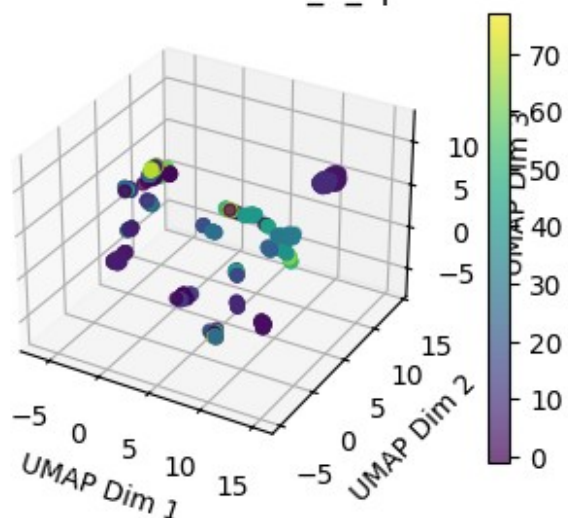
pol\_X\_kpca Silhouette Score (Ignoring Noise): 0.97337943  
clusters: 2

3D Visualization pol\_X\_kpca



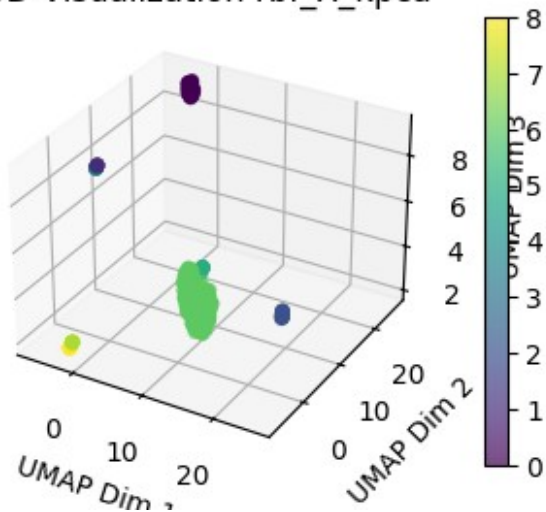
cos\_X\_kpca Silhouette Score (Ignoring Noise): 0.555738  
clusters: 79

3D Visualization cos\_X\_kpca



rbf\_X\_kpca Silhouette Score (Ignoring Noise): 0.31911534  
clusters: 9

3D Visualization rbf\_X\_kpca



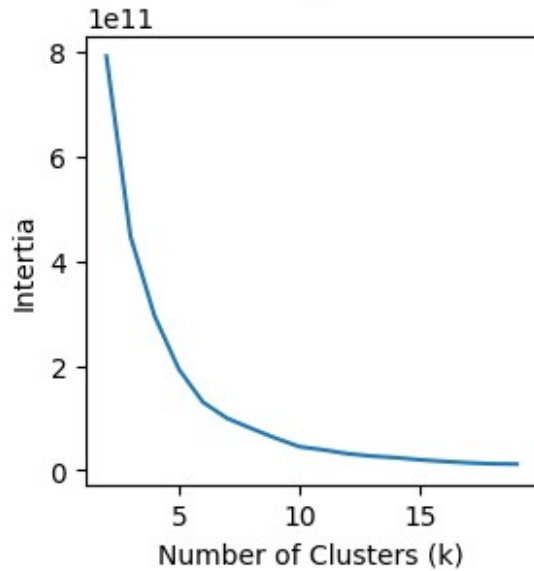
#### Observation

- Non-reduced with kpca poly kernel has highest scoring but it diverts data into 2 clusters only.
- Non-reduced with kpca rbf kernel has high scoring and also gives 11 different clusters.

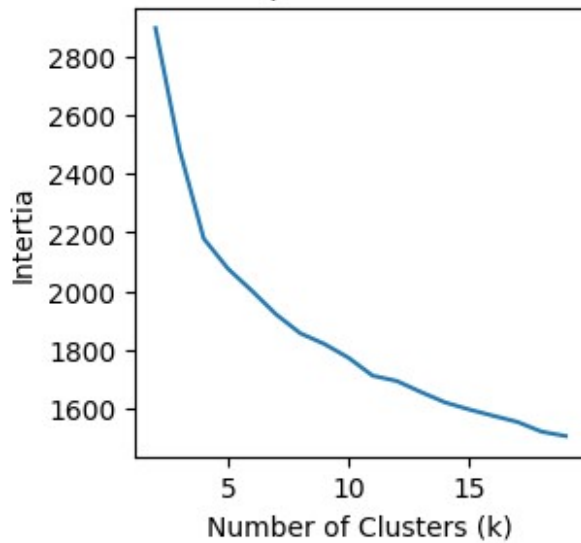
```
# KMeans Clustering algorithm use
plot_kmeans(df_final, 'df_final')
plot_kmeans(embedding_matrix, 'embedding_matrix')
plot_kmeans(X_scaled, 'X_scaled')
plot_kmeans(X, 'X')
plot_kmeans(pol_X_kpca, 'pol_X_kpca')
```

```
plot_kmeans(cos_X_kpca, 'cos_X_kpca')  
plot_kmeans(rbf_X_kpca, 'rbf_X_kpca')
```

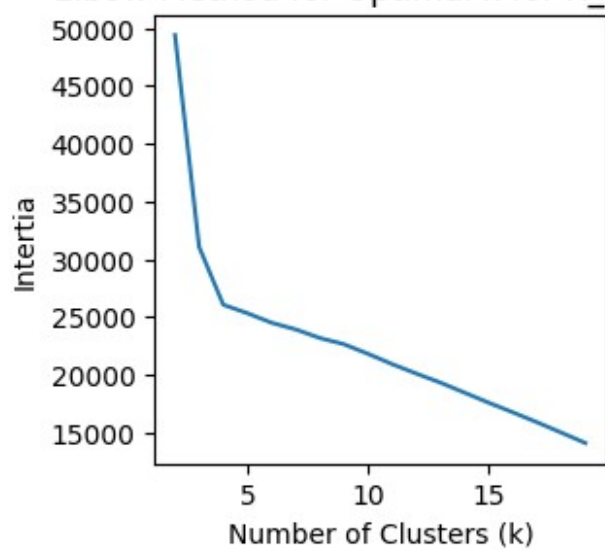
Elbow Method for Optimal k for df\_final



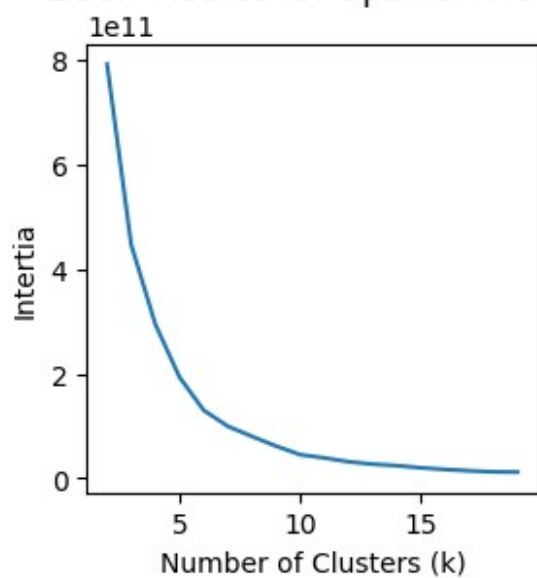
Elbow Method for Optimal k for embedding\_matrix



Elbow Method for Optimal k for X\_scaled

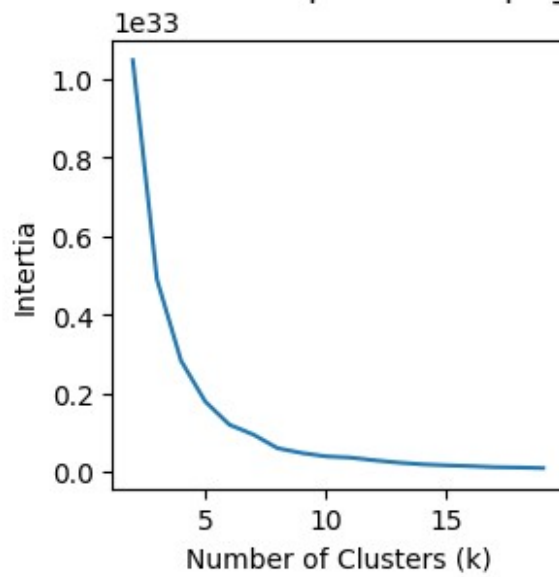


Elbow Method for Optimal k for X

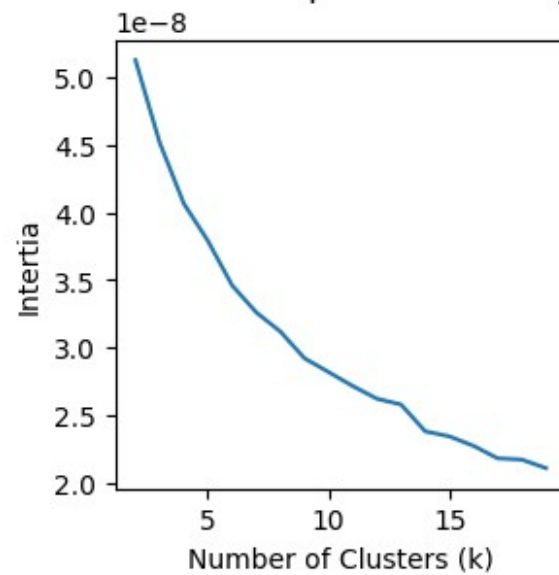




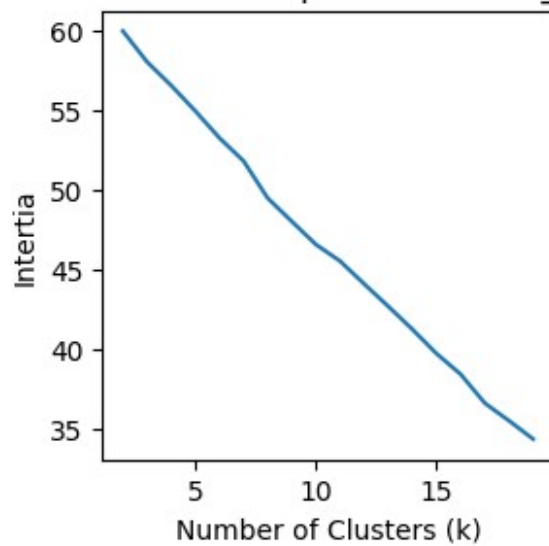
Elbow Method for Optimal k for pol\_X\_kpca



Elbow Method for Optimal k for cos\_X\_kpca



Elbow Method for Optimal k for rbf\_X\_kpca

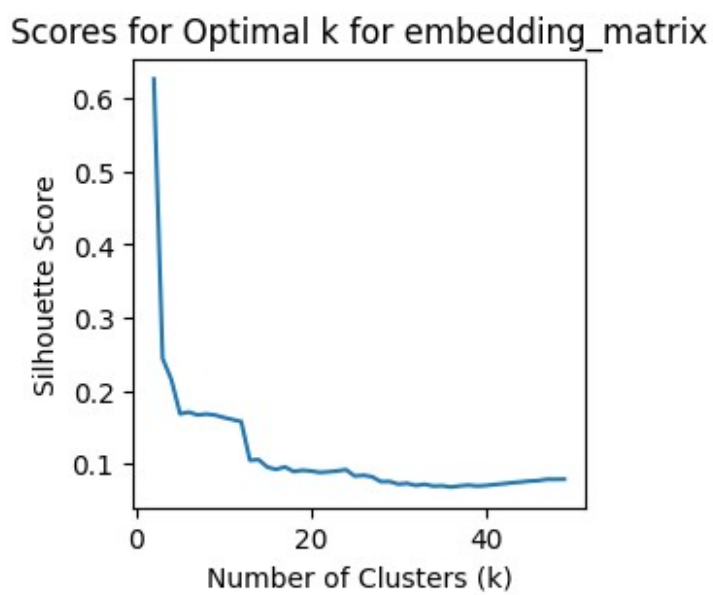
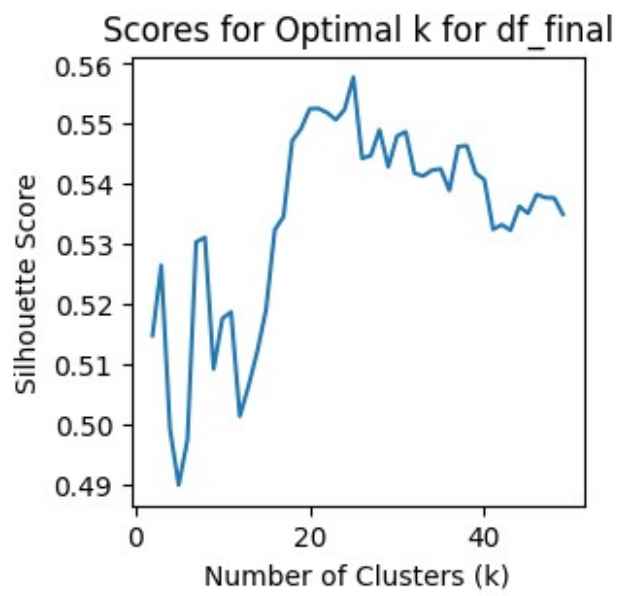


#### Observation

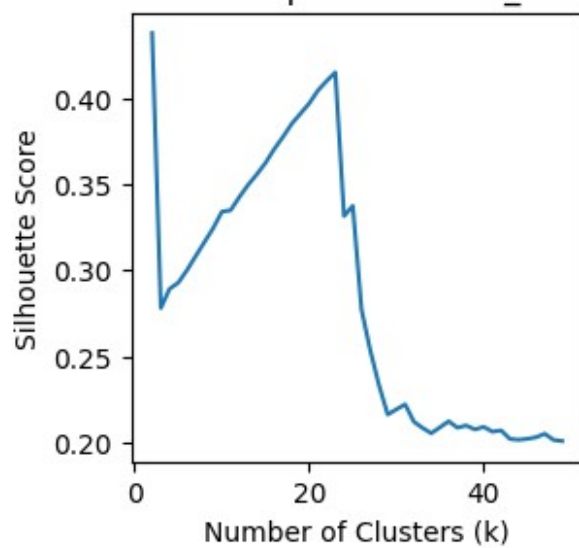
- With KMeans rbf doesn't go well maybe need to increase number of clusters.
- Mostly No elbow or very fine Elbow point is discovered with all type of data.

*# Agglomerative Clustering Plotting with different scores and n\_values.*

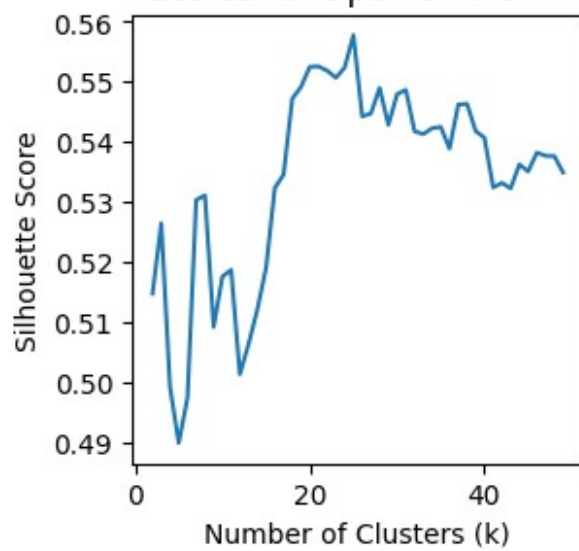
```
AgglomerativeClustering_plot(df_final, 'df_final')
AgglomerativeClustering_plot(embedding_matrix, 'embedding_matrix')
AgglomerativeClustering_plot(X_scaled, 'X_scaled')
AgglomerativeClustering_plot(X, 'X')
AgglomerativeClustering_plot(pol_X_kpca, 'pol_X_kpca')
AgglomerativeClustering_plot(cos_X_kpca, 'cos_X_kpca')
AgglomerativeClustering_plot(rbf_X_kpca, 'rbf_X_kpca')
```



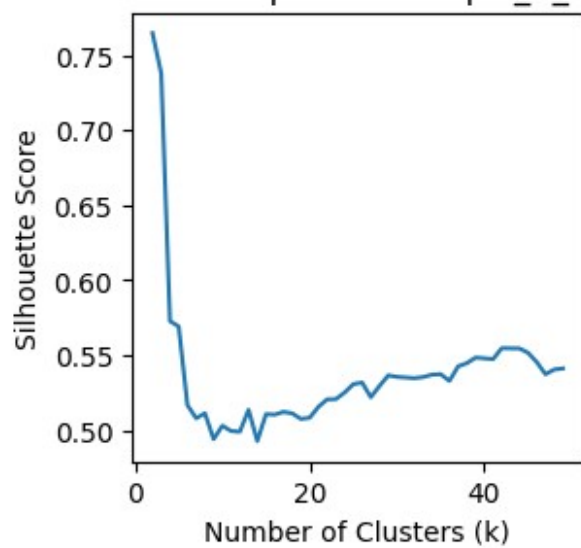
Scores for Optimal k for X\_scaled



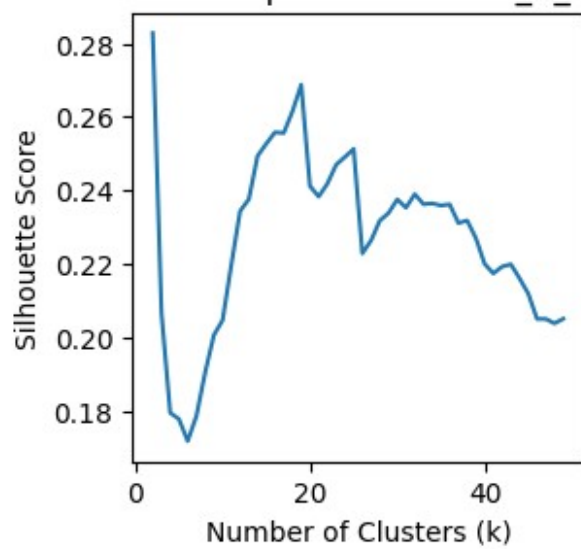
Scores for Optimal k for X

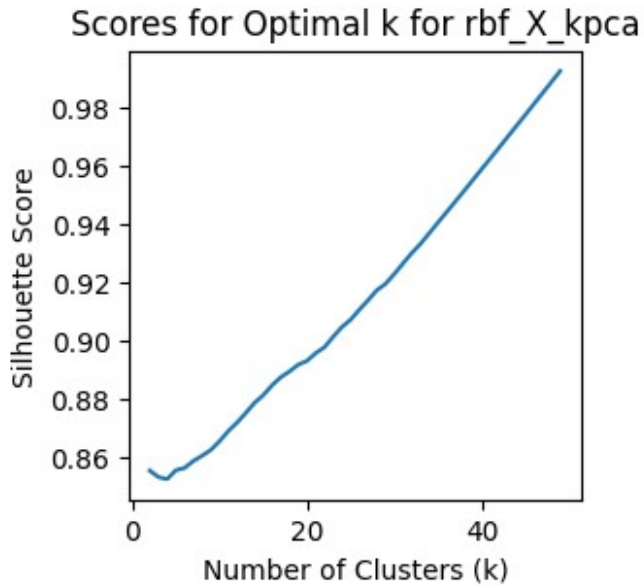


Scores for Optimal k for pol\_X\_kpca



Scores for Optimal k for cos\_X\_kpca





#### Observation

- Again only kpca with kernel **rbf** gives a better score but still the score only goes up with more clustering.
- Reduced dataset with no scaling and non-reduced dataset also worked fine with clustering scores
- X dataset best\_n = 23-28
- final\_df best\_n = 25-28

```
DBSCAN_train(df_final,'df_final')
DBSCAN_train(embedding_matrix,'embedding_matrix')
DBSCAN_train(X_scaled,'X_scaled')
DBSCAN_train(X,'X')
DBSCAN_train(pol_X_kpca,'pol_X_kpca')
DBSCAN_train(cos_X_kpca,'cos_X_kpca')
DBSCAN_train(rbf_X_kpca,'rbf_X_kpca')
```

```
embedding_matrix Best DBSCAN Params: eps=0.8, min_samples=10,
silhouette_score=0.402, Cluster:2
```

```
rbf_X_kpca Best DBSCAN Params: eps=0.05, min_samples=5,
silhouette_score=0.857, Cluster:2
```

#### Observation

- We have again much better results with rbf kernel of kpca, but this time DBSCAN couldn't get the non-linear relationship

## Key-Findings, Walkthrough, Summary

- There could be multiple best clusters approach based on need and more data.
- HDBSCAN is best clustering algorithm for this data

- Worked with multiple combination of data pre-processing such as PCA+TruncatedSVD+UMAP, KPCA+UMAP, KPCA with different kernels.
- I used different techniques to catch different features uniquely such as semantic as 1, Numeric as 1, and Categorical as 1, also tried if combination of all would work better or not.
- Worked with multiple clustering algorithm to get the best one out of it, as data has non-linearity I have used DBSCAN, HDBSCAN as it also has categorical data so used tree based algorithm AgglomerativeClustering.
- To work with non-linearity and Kmeans I also used KPCA, UMAP.
- To map in 2D I used MDS, and TSNE like techniques.
- PLotted graph with scores and clusters to understanding visually and have a better decision making.
- HDBSCAN with KPCA and UMAP works well as KPCA and UMAP preserves both clustering and semantic relations and HDBSCAN works well with higher density making it a perfect model for our dataset.

## Best Model To Choose

- Using **HDBSCAN with UMAP reduction technique** give the best clustering result with **11 clusters** and **score = 89.6**

## Future Recommendation and sugesstion

- Could fine tune HDBSCAN, DBSCAN algorithms for better results.
- using KPCA with reduced data using different approaches with HDBSCAN for getting more better clustering results.
- Different distance metrics and scoring algorithm could be used.
- Using Similarity Matrix with these for Song Recommendation System.